

New ways of accessing data: Parquet format and data in the cloud

Lino Galiana

Version of September 27, 2026

Parquet is a very popular data format, thanks to its technical characteristics (column orientation, compression, interoperability...), its open-source nature and the rich associated ecosystem, whose most prominent frameworks are **Arrow** and **DuckDB**. Added to these numerous advantages are native integration with cloud infrastructures based on **S3**, numerous extensions for handling complex data such as geographic data or, more recently, the WASM porting of **DuckDB** enabling responsive applications involving data transformations to be built directly from the browser.

 **Automatically generated PDF.** This document is produced from the course website and may contain formatting issues. For an up-to-date version adapted to your reading, visit the course page: pythonds.linogaliana.fr.

Table of contents

1 Contextual Elements	2
1.1 Principles of Data Storage	2
1.2 File-Based Storage	2
1.3 Storing Data in a Database	4
2 The Parquet format	5
2.1 Limitations of the CSV format	5
2.2 The Rise of the Parquet Format	6
2.3 A format designed for analysis — not just big data	7
2.4 Reading a Parquet file in Python : example	8
2.5 Arrow	8
2.6 DuckDB	9
2.7 Exercises to Learn More	10
3 Data in the Cloud	18
3.1 Why Not Use Dropbox or Google Drive ?	18
3.2 The S3 System	18
3.3 How to do it with Python?	19
3.4 Practical Case: Storing Your Project's Data on SSP Cloud	19
3.5 To go further	23
Informations additionnelles	23
Bibliography	25

In the previous chapters, we explored how to retrieve and harmonize data from various source: CSV files, APIs, web scraping, and more. However, no overview of data access methods would be complete without mentioning a relatively recent addition to the data ecosystem: the **Parquet** data format.

Thanks to its technical advantages - specifically designed for analytical workloads - and its seamless integration with Python, **Parquet** is becoming increasingly essential. It has, in fact, become a key component of modern cloud infrastructures, which have emerged since the mid-2010s as the standard environment for data science workflows.¹

¹For more details, see the [production deployment course](#) by Romain Avouac and myself. Apologies to non-French-speaking readers—this resource is currently only available in French.

💡 Objectives

- Understand the challenges involved in storing and processing different types of data formats
- Distinguish between file-based storage and database systems
- Discover **Parquet** format and its advantages over flat files and proprietary formats
- Learn how to work with **Parquet** data using **Arrow** and **DuckDB**
- Explore the implications of cloud-based storage and how **Python** can adapt to modern data infrastructures

1 Contextual Elements

1.1 Principles of Data Storage

Before exploring the advantages of the **Parquet** format, it is helpful to briefly review how data is stored and made accessible to a processing language like **Python**².

Two main approaches coexist: **file-based storage** and **relational database storage**. The key distinction between these paradigms lies in how access to data is structured and managed.

1.2 File-Based Storage

1.2.a Flat Files

In a flat file, data is organized in a linear fashion, with values typically separated by a delimiter such as a comma, semicolon, or tab. Here's a simple example using a **.csv** file:

```
nom ; profession
Astérix ;
Obélix ; Tailleur de menhir ;
Assurancetourix ; Barde
```

Python can easily structure this information:

	nom	profession
0	Astérix	
1	Obélix	Tailleur de menhir
2	Assurancetourix	Barde

1. **StringIO** lets us treat the character string as the content of a file.

These are referred to as **flat files** because all the records related to a given observation are stored sequentially, without any hierarchical structure.

1.2.b Hierarchical files

Other formats, such as **JSON**, structure data in a hierarchical way:

```
[
{
```

²For a deeper discussion of data format selection challenges, see A. Dondon and P. Lamarche [1].

```
[{"nom": "Astérix"}, {"nom": "Obélix", "profession": "Tailleur de menhir"}, {"nom": "Assurancetourix", "profession": "Barde"}]
```

In such cases, when information is missing - as in the line for “Astérix” - you don’t see two delimiters side by side. Instead, the missing data is simply left out.

Caution

The difference between a `.csv` file and a `JSON` file lies not only in their format but also in their underlying logic of data storage.

`JSON` is a non-tabular format that offers greater flexibility: it allows the structure of the data to evolve over time without requiring previous entries to be modified or recompiled. This makes it particularly well-suited for dynamic data collection environments, such as APIs.

For instance, a website that begins collecting a new piece of information doesn’t need to retroactively update all prior records. It can simply add the new field to the relevant entries—omitting it from others where the information wasn’t collected.

It is then up to the query tool—such as `Python` or another platform—to reconcile and link data across these differing structures.

This flexible approach underpins many `NoSQL` databases (like `ElasticSearch`), which play a central role in the *big data* ecosystem.

1.2.c Data split across multiple files

It is common for a single observation to be distributed across multiple files in different formats. In geomatics, for example, geographic boundaries are often stored separately from the data that gives them context:

- In some cases, everything is bundled into a single file that contains both the geographic shapes and the associated attribute values. This is the approach used by formats like `GeoJSON`;
- In other cases, the data is split across several files, and reading the full information - geographic shapes, data for each zone, projection system, etc. - requires combining them. This is the case with the `Shapefile` format.

When data is distributed across multiple files, this is up to the processing tool (e.g., `Python`) to perform the necessary joins and link the information together.

1.2.d The role of the file system

The **file system** enables the computer to locate files physically on the disk. It plays a central role in file management, handling file naming, folder hierarchy, and access permissions.

1.3 Storing Data in a Database

The logic behind databases differs fundamentally and is more systematic in nature. A **relational database** is managed by a **Database Management System (DBMS)**, which provides capabilities for:

- storing coherent datasets;
- performing updates (insertions, deletions, modifications);
- controlling access (user permissions, query types, etc.).

Data is organized into **tables** connected by **relationships**, often structured according to a **star schema**:




Figure 1: Source: [Databricks Documentation](#)

The software managing the database links these tables together using **SQL** queries. One of the most powerful and widely used systems for this purpose is **PostgreSQL**.

Python can interact with databases by issuing SQL queries. Historically, packages such as **sqlalchemy** and **psycopg2** have been standard tools for communicating with **PostgreSQL** databases, enabling both reading and updating operations. More recently, **DuckDB** has emerged as a lightweight and user-friendly alternative for querying relational data, and it will be discussed again in the context of the **Parquet** format.

Why File-Based Storage is gaining popularity

The increasing popularity of file-based storage in the data science ecosystem is due to a number of technical and practical advantages that make it well-suited for modern analytical workflows.

Files are generally more lightweight and easier to manage than databases. They do not require the installation or maintenance of specialized software; a basic file system, available on every operating system, is sufficient for accessing them.

Reading a file in **Python** simply involves using a library such as **Pandas**. In contrast, interacting with a database typically requires:

- installing and configuring a DBMS (e.g., **PostgreSQL**, **MySQL**);
- managing network connections;
- relying on libraries such as **sqlalchemy** or **psycopg2**.

This additional complexity makes file-based workflows more flexible and faster for exploratory tasks.

However, this simplicity comes with limitations. File-based systems generally lack fine-grained access control. For instance, preventing a user from modifying or deleting a file is difficult without duplicating it and working from a copy. This represents a constraint in multi-user environments, although cloud-based storage solutions — particularly **S3** technology, which will be addressed later — offer effective remedies.

The primary reason why files are often preferred over DBMSs lies in the nature of the operations being performed. Relational databases are particularly well-suited for contexts involving frequent updates or complex operations on structured data — a typical **application logic**, in which data is continuously evolving (through insertions, updates, or deletions).

By contrast, in **analytical** contexts, the focus is on reading and temporarily manipulating data without altering the original source. The objective is to query, aggregate, and filter — not to persist changes. In such scenarios, files (especially when stored in optimized formats like **Parquet**) are ideal: they offer fast read times, high portability, and eliminate the overhead associated with running a full database engine.

2 The **Parquet** format

The **CSV** format has long been popular for its simplicity:

- It is **human-readable** (any text editor can open it);
- It relies on a **simple tabular structure**, well-suited for many analytical situations;
- It is **universal** and interoperable, as it is not tied to any specific software.

However, this simplicity comes at a cost. Several limitations of the **CSV** format have led to the emergence of more efficient formats for data analysis, such as **Parquet**.

2.1 Limitations of the **CSV** format

CSV is a **heavy** format:

- It is not compressed, increasing its disk size;
- All data is stored as raw text. Data type optimization (integer, float, string, etc.) is deferred to the importing library (like **Pandas**), which must **scan the data** at load time—slowing down performance and increasing error risk.

CSV is **row-oriented**:

- To access a specific column, every **row must be read**, and the relevant column extracted;
- This model performs poorly when only a **subset of columns** is needed—a common case in data science.

CSV is **expensive to modify**:

- Adding a column or inserting intermediate data requires **rewriting the entire file**. For example, adding a **hair** column would mean generating a new version of the file:

```
name ; hair ; profession
Asterix ; blond ;
Obelix ; redhead ; Menhir sculptor
Assurancetourix ; blond ; Bard
```

i About proprietary formats

Most data science tools offer their own serialization formats:

- `.pickle` for `Python`,
- `.rda` or `.RData` for `R`,
- `.dta` for `Stata`,
- `.sas7bdat` for `SAS`.

However, these formats are **proprietary** or **tightly coupled to one language**, which raises interoperability issues. For example, `Python` cannot natively read `.sas7bdat`. Even when third-party libraries exist, the lack of official documentation makes support unreliable.

In this regard, despite its limitations, `.csv` remains popular for its universality. But the `Parquet` format offers both this portability and significantly better performance.

2.2 The Rise of the `Parquet` Format

To address these limitations, the `Parquet` format—developed as an [Apache open-source project](#) offers a radically different approach.

Its key characteristic: it is **column-oriented**. Unlike CSVs, data for each column is stored **separately**. This allows:

- loading only the columns relevant for analysis;
- more efficient data compression;
- significantly faster selective queries.

Here's a diagram from the [Upsolver blog](#) illustrating the difference between row-based and columnar storage:




Figure 2: Parquet vs CSV

In our example, you could read the `profession` column without parsing names, making access faster (ignore the `pyarrow.Table` element—we'll return to it later):

```
pyarrow.Table
nom : string
profession: string
----
nom : ["Astérix ", "Obélix ", "Assurancetourix "]
profession: [["", "Tailleur de menhir", "Barde"]]
```

Thanks to the column-oriented structure, it is possible to read only a single variable (such as `profession`) without having to scan every row in the file.

```
path
└─ to
   └─ table
      ├── gender=male
      │   ├── country=US
      │   │   └─ data.parquet
      │   ├── country=CN
      │   │   └─ data.parquet
      └── gender=female
          ├── country=US
          │   └─ data.parquet
          ├── country=CN
          │   └─ data.parquet
```

When read, the entire dataset is reconstructed into a tabular format:

```
root
|-- name: string (nullable = true)
|-- age: long (nullable = true)
|-- gender: string (nullable = true)
|-- country: string (nullable = true)
```

2.3 A format designed for analysis — not just big data

As emphasized in the Upsolver blog:

Complex data such as logs and event streams would need to be represented as a table with hundreds or thousands of columns, and many millions of rows. Storing this table in a row-based format such as CSV would mean:

- *Queries will take longer to run since more data needs to be scanned...*
- *Storage will be more costly since CSVs are not compressed as efficiently as Parquet*

However, the `Parquet` format is **not limited to *big data* architectures**. Its advantages are accessible to anyone who produces or works with datasets, regardless of scale:

- significantly reduced file sizes;
- fast, reliable, and memory-efficient data import.

The translation of this chapter is still in progress. The remaining sections will be available soon.

2.4 Reading a Parquet file in Python: example

There are many libraries that integrate well with the Parquet format, but the two most important to know are PyArrow and DuckDB. These libraries were previously mentioned as alternatives to Pandas for handling larger-than-memory datasets. They are often used to perform heavy initial operations before converting the results into a lightweight `pd.DataFrame` for further analysis.

The PyArrow library enables efficient reading and writing of Parquet files by taking advantage of their columnar structure³. It operates on a `pyarrow.Table` object, which—after processing—can be converted to a Pandas `DataFrame` to leverage the broader capabilities of the Pandas ecosystem.

The DuckDB library allows you to query Parquet files directly using SQL, without loading the entire file into memory. In essence, it brings the database philosophy (structured queries via SQL) to file-based storage. The results of such queries can also be converted into a Pandas `DataFrame`, combining the convenience of Pandas with the efficiency of an embedded SQL engine.

A lesser-known but valuable feature of DuckDB is its ability to perform SQL queries directly on a Pandas `DataFrame`. This can be especially useful when Pandas syntax becomes verbose or cumbersome—for example, when computing a new column based on grouped statistics, where SQL expressions can be more concise and readable.

Tip

Using `pa` for `pyarrow` and `pq` for `pyarrow.parquet` is a widely adopted convention, much like using `pd` for `pandas`.

To demonstrate these features, we will use a dataset derived from the synthetic census data published by Insee, the French national statistics agency.

```
import requests
import pyarrow.parquet as pq

# Example Parquet
url = "https://minio.lab.sspcloud.fr/projet-formation/bonnes-pratiques/data/RPindivodus/REGION=93/part-0.parquet"

# Télécharger le fichier et l'enregistrer en local
with open("example.parquet", "wb") as f:
    response = requests.get(url)
    f.write(response.content)
```

2.5 Arrow

To fully benefit from the optimizations provided by the Parquet format, it is recommended to use `pyarrow.dataset`. This approach makes it possible to take full advantage of the performance gains offered by the combination of Parquet and Arrow, which are not always accessible when reading Parquet files using other methods available in the Arrow ecosystem (as will be explored in upcoming exercises).

```
import pyarrow.dataset as ds

dataset = ds.dataset(
    "example.parquet"
```

³It is recommended to regularly consult official `pyarrow` documentation on [reading and writing files and data manipulation](#).


```
.scanner(columns = ["AGED", "IPONDI", "DEPT"])
table = dataset.to_table()
table
```

```
pyarrow.Table
AGED: int32
IPONDI: double
DEPT: dictionary<values=string, indices=int32, ordered=0>
----
AGED:          [[9,12,40,70,52, ..., 29,66,72,75,77],[46,76,46,32,2, ..., 7,5,37,29,4], ...,
[67,37,45,56,75, ..., 64,37,47,20,18],[16,25,51,6,11, ..., 93,90,92,21,65]]
IPONDI: [[2.73018871840726,2.73018871840726,2.73018871840726,0.954760150327854,3.75907197064638, ..., 3.27143319621654,
[3.02627578376137,3.01215358930406,3.01215358930406,2.93136309038958,2.93136309038958, ..., 2.96848755763453,2.96848755
[2.57931132917563,2.85579410739065,0.845993555838931,2.50296716736141,3.70786113613679, ..., 3.08375347880892,2.8803880
[3.22776230929947,3.22776230929947,3.22776230929947,3.29380242174036,3.29380242174036, ..., 5.00000768518755,5.00000768
DEPT: [ -- dictionary:
["01","02","03","04","05", ..., "95","971","972","973","974"] -- indices:
[5,5,5,5,5, ..., 5,5,5,5,5], -- dictionary:
["01","02","03","04","05", ..., "95","971","972","973","974"] -- indices:
[5,5,5,5,5, ..., 5,5,5,5,5], ..., -- dictionary:
["01","02","03","04","05", ..., "95","971","972","973","974"] -- indices:
[84,84,84,84,84, ..., 84,84,84,84,84], -- dictionary:
["01","02","03","04","05", ..., "95","971","972","973","974"] -- indices:
[84,84,84,84,84, ..., 84,84,84,84,84]]
```

To import and process these data, one can either keep the data in `pyarrow.Table` format or convert it into a `pandas.DataFrame`. The second option is slower but has the advantage of enabling all the manipulations offered by the `pandas` ecosystem, which is generally more familiar than that of `Arrow`.

2.6 DuckDB

```
import duckdb
duckdb.sql("""
FROM read_parquet('example.parquet')
SELECT AGED, IPONDI, DEPT
""")
```

AGED	IPONDI	DEPT
int32	double	varchar
9	2.73018871840726	06
12	2.73018871840726	06
40	2.73018871840726	06
70	0.954760150327854	06
52	3.75907197064638	06
82	3.21622922493506	06
6	3.44170061276923	06
12	3.44170061276923	06
15	3.44170061276923	06
43	3.44170061276923	06
.	.	.
.	.	.
.	.	.

68	2.73018871840726	06
35	3.46310256220757	06
2	3.46310256220757	06
37	3.46310256220757	06
84	3.69787960424482	06
81	4.7717265388427	06
81	4.7717265388427	06
51	3.60566450823737	06
25	3.60566450823737	06
13	3.60566450823737	06

? rows 3 columns
(>9999 rows, 20 shown)

2.7 Exercises to Learn More

The following is a series of exercises adapted from the [production deployment of data science projects](#) course, which Romain Avouac and I teach in the final year of the engineering program at ENSAE.

These exercises progressively illustrate some of the key concepts discussed above, while also emphasizing best practices for working with large-scale data. Solutions to all exercises are available on the corresponding course page.

In this practical section, we will explore how to use the **Parquet** format as efficiently as possible. To compare different data formats and access strategies, we will **measure and compare the execution time and memory usage of a standard query**. We will begin with a lightweight example that compares the performance of reading data in **CSV** versus **Parquet** format.

To do so, we will first retrieve a dataset in **Parquet** format. We suggest using the detailed and anonymized French population census data, which contains approximately 20 million rows and 80 columns. The code to download and access this dataset is provided below.

```
import pyarrow.parquet as pq
import pyarrow as pa
import os

# Définir le fichier de destination
filename_table_individu = "data/RPindividus.parquet"

# Copier le fichier depuis le stockage distant (remplacer par une méthode adaptée si nécessaire)
os.system("mc cp s3/projet-formation/bonnes-pratiques/data/RPindividus.parquet data/RPindividus.parquet")

# Charger le fichier Parquet
table = pq.read_table(filename_table_individu)
df = table.to_pandas()

# Filtrer les données pour REGION == "24"
df_filtered = df.loc[df["REGION"] == "24"]

# Sauvegarder en CSV
df_filtered.to_csv("data/RPindividus_24.csv", index=False)

# Sauvegarder en Parquet
pq.write_table(pa.Table.from_pandas(df_filtered), "data/RPindividus_24.parquet")
```

Line 9

This line of code uses the Minio Client utility available on the [SSPCloud](#). If you're not on this infrastructure, please refer to the dedicated box.

! If you are not on SSPCloud

You will need to replace the line

```
os.system("mc cp s3/projet-formation/bonnes-pratiques/data/RPindividus.parquet data/
RPindividus.parquet")
```

which uses the `mc` command-line tool, with code that downloads this data from the URL <https://projet-formation.minio.lab.sspcloud.fr/bonnes-pratiques/data/RPindividus.parquet>.

There are many ways to do this. For instance, you can use plain `Python` with `requests`. If you have `curl` installed, you can use it as well. Via `Python`, this would translate to the command: `os.system("curl -o data/RPindividus.parquet https://projet-formation/bonnes-pratiques/data/RPindividus.parquet")`.

These exercises will make use of `Python` decorators - functions that modify or extend the behavior of other functions. In our case, we will define a function that performs a series of operations, and then apply a decorator to it that tracks both memory usage and execution time.

💡 Part 1: From CSV to Parquet

- Create a notebook named `benchmark_parquet.ipynb` to perform various performance comparisons throughout the application.
- Define a custom decorator that will be used to benchmark the `Python` code by measuring execution time and memory usage.

Click to expand and view the code for the performance-measuring decorator. `::: {cell execution_count=12} `` {python .cell-code} import time from memory_profiler import memory_usage from functools import wraps import warnings`

```
def convert_size(size_bytes):
    if size_bytes == 0:
        return "0B"
    size_name = ("B", "KB", "MB", "GB", "TB", "PB", "EB", "ZB", "YB")
    i = int(math.floor(math.log(size_bytes, 1024)))
    p = math.pow(1024, i)
    s = round(size_bytes / p, 2)
    return "%s %s" % (s, size_name[i])

# Decorator to measure execution time and memory usage
def measure_performance(func, return_output=False):
    @wraps(func)
    def wrapper(return_output=False, *args, **kwargs):
        warnings.filterwarnings("ignore")
        start_time = time.time()
        mem_usage = memory_usage((func, args, kwargs), interval=0.1)
        end_time = time.time()
        warnings.filterwarnings("always")

        exec_time = end_time - start_time
        peak_mem = max(mem_usage) # Peak memory usage
        exec_time_formatted = f"\033[92m{exec_time:.4f} sec\033[0m"
```

```

        peak_mem_formatted = f"\033[92m{convert_size(1024*peak_mem)}\033[0m"

        print(f"{func.__name__} - Execution Time: {exec_time_formatted} | Peak Memory Usage: {peak_mem_formatted}")
        if return_output is True:
            return func(*args, **kwargs)

    return wrapper

```

```

:::

</details>

* Reuse this code to wrap the logic for constructing the age pyramid into a function named `process_csv_appli1`.

<details>

<summary>
Click to expand and view the code used to measure the performance of reading CSV files.
</summary>

::: {.cell execution_count=13}
``` {.python .cell-code}
Apply the decorator to functions
@measure_performance
def process_csv_appli1(*args, **kwargs):
 df = pd.read_csv("data/RPindivividus_24.csv")
 return (
 df.loc[df["DEPT"] == 36]
 .groupby(["AGED", "DEPT"])["IPONDI"]
 .sum().reset_index()
 .rename(columns={"IPONDI": "n_indiv"})
)

```

```

:::

```

- Run `process_csv_appli1()` and `process_csv_appli1(return_output=True)` to observe performance and optionally return the processed data.
- Using the same approach, define a new function named `process_parquet_appli1`, this time based on the `data/RPindivividus_24.parquet` file, and load it using Pandas' `read_parquet` function.
- Compare the performance (execution time and memory usage) of the two methods using the benchmarking decorator.

#### Full correction

```

import math
import pandas as pd
import time
from memory_profiler import memory_usage
from functools import wraps
import warnings

def convert_size(size_bytes):

```

```

if size_bytes == 0:
 return "0B"
size_name = ("B", "KB", "MB", "GB", "TB", "PB", "EB", "ZB", "YB")
i = int(math.floor(math.log(size_bytes, 1024)))
p = math.pow(1024, i)
s = round(size_bytes / p, 2)
return "%s %s" % (s, size_name[i])

Decorator to measure execution time and memory usage
def measure_performance(func, return_output=False):
 @wraps(func)
 def wrapper(return_output=False, *args, **kwargs):
 warnings.filterwarnings("ignore")
 start_time = time.time()
 mem_usage = memory_usage((func, args, kwargs), interval=0.1)
 end_time = time.time()
 warnings.filterwarnings("always")

 exec_time = end_time - start_time
 peak_mem = max(mem_usage) # Peak memory usage
 exec_time_formatted = f"\033[92m{exec_time:.4f} sec\033[0m"
 peak_mem_formatted = f"\033[92m{convert_size(1024*peak_mem)}\033[0m"

 print(f"{func.__name__} - Execution Time: {exec_time_formatted} | Peak Memory Usage: {peak_mem_formatted}")
 if return_output is True:
 return func(*args, **kwargs)

 return wrapper

Apply the decorator to functions
@measure_performance
def process_csv(*args, **kwargs):
 df = pd.read_csv("data/RPindivividus_24.csv")
 return (
 df.loc[df["DEPT"] == 36]
 .groupby(["AGED", "DEPT"])["IPONDI"]
 .sum().reset_index()
 .rename(columns={"IPONDI": "n_indiv"})
)

@measure_performance
def process_parquet(*args, **kwargs):
 df = pd.read_parquet("data/RPindivividus_24.parquet")
 return (
 df.loc[df["DEPT"] == 36]
 .groupby(["AGED", "DEPT"])["IPONDI"]
 .sum().reset_index()
 .rename(columns={"IPONDI": "n_indiv"})
)

process_csv()
process_parquet()

```

? What seems to be the limitation of the `read_parquet` function?

Although we already observe a significant speed improvement during file reading, we are not fully leveraging the optimizations provided by the `Parquet` format. This is because the data is immediately loaded into a `Pandas DataFrame`, where transformations are applied afterward.

As a result, we miss out on one of `Parquet`'s core performance features: **predicate pushdown**. This optimization allows filters to be applied as early as possible—at the file scan level—so that only the relevant columns and rows are read into memory. By bypassing this mechanism, we lose much of what makes `Parquet` so efficient in analytical workflows.

## 💡 Part 2: Leveraging Lazy Evaluation and the Optimizations of Arrow or DuckDB

In the previous section, we observed a **significant improvement in read times** when switching from `CSV` to `Parquet`. However, **memory usage remained high**, even though only a small portion of the data was actually used.

In this section, we will explore how to take advantage of **lazy evaluation** and **execution plan optimizations** offered by `Arrow` to fully unlock the performance benefits of the `Parquet` format.

- Open the file `data/RPindividus_24.parquet` using `pyarrow.dataset`. Check the class of the resulting object.
- Run the code below to read a sample of the data:

```
(
 dataset.scanner()
 .head(5)
 .to_pandas()
)
```

Can you identify the difference compared to the previous approach? Consult the documentation for the `to_table` method—do you understand what it does and why it matters?

- Create a function `summarize_parquet_arrow` (and a corresponding `summarize_parquet_duckdb`) that imports the data using `pyarrow.dataset` (or `DuckDB`) and performs the required aggregation.
- Use the benchmarking decorator to compare the performance (execution time and memory usage) of the three approaches: reading and processing `Parquet` data using `Pandas`, `PyArrow`, and `DuckDB`.

Correction

```
import duckdb
import pyarrow.dataset as ds

@measure_performance
def summarize_parquet_duckdb(*args, **kwargs):
 con = duckdb.connect(":memory:")
 query = """
 FROM read_parquet('data/RPindividus_24.parquet')
 SELECT AGED, DEPT, SUM(IPONDI) AS n_indiv
 GROUP BY AGED, DEPT
 """

 return (con.sql(query).to_df())

@measure_performance
def summarize_parquet_arrow(*args, **kwargs):

 dataset = ds.dataset("data/RPindividus_24.parquet", format="parquet")
```

```

table = dataset.to_table()
grouped_table = (
 table
 .group_by(["AGED", "DEPT"])
 .aggregate(["IPONDI", "sum"])
 .rename_columns(["AGED", "DEPT", "n_indiv"])
 .to_pandas()
)


return (
 grouped_table
)

process_parquet()
summarize_parquet_duckdb()
summarize_parquet_arrow()

```

With lazy evaluation, the process unfolds in several stages:

- **Arrow** or **DuckDB** receives a set of instructions, builds an execution plan, optimizes it, and then executes the query;
- Only the final result of this pipeline is returned to **Python**, rather than the entire dataset.



### 💡 Part 3a: What Happens When We Filter Rows?

Let us now add a row-level filtering step to our queries:

- With **DuckDB**, modify the SQL query to include a **WHERE** clause:  
`WHERE DEPT IN ('18', '28', '36')`
- With **Arrow**, update the `to_table` call as follows:  
`dataset.to_table(filter=pc.field("DEPT").isin(['18', '28', '36']))`

```

import pyarrow.dataset as ds
import pyarrow.compute as pc
import duckdb

@measure_performance
def summarize_filter_parquet_arrow(*args, **kwargs):

 dataset = ds.dataset("data/RPindividus.parquet", format="parquet")
 table = dataset.to_table(filter=pc.field("DEPT").isin(['18', '28', '36']))
 grouped_table = (
 table
 .group_by(["AGED", "DEPT"])
 .aggregate(["IPONDI", "sum"])
 .rename_columns(["AGED", "DEPT", "n_indiv"])
 .to_pandas()
)

```

```

 return (
 grouped_table
)

@measure_performance
def summarize_filter_parquet_duckdb(*args, **kwargs):
 con = duckdb.connect(":memory:")
 query = """
 FROM read_parquet('data/RPindividus_24.parquet')
 SELECT AGED, DEPT, SUM(IPONDI) AS n_indiv
 WHERE DEPT IN ('11','31','34')
 GROUP BY AGED, DEPT
 """

 return (con.sql(query).to_df())

summarize_filter_parquet_arrow()
summarize_filter_parquet_duckdb()

```

**?** Why do row filters not improve performance (and sometimes even slow things down), unlike column filters?

This is because data is not stored in row blocks the way it is in column blocks. As a result, filtering rows does not allow the system to skip over large sections of the file as efficiently.

Fortunately, there is a solution: **partitioning**.

### 💡 Part 3: Partitioned Parquet

*Lazy evaluation* and the optimizations available through **Arrow** already provide significant performance improvements. But we can go even further. When you know in advance that your queries will frequently **filter data based on a specific variable**, it is highly advantageous to **partition** the **Parquet** file using that variable.

1. Review the documentation for `pyarrow.parquet.write_to_dataset` to understand how to define a partitioning key when writing a **Parquet** file. Several approaches are available.
2. Import the full individuals table from the census using `pyarrow.dataset.dataset("data/RPindividus.parquet")`, and export it as a partitioned dataset to `"data/RPindividus_partitionne.parquet"`, using both **REGION** and **DEPT** as partitioning keys.
3. Explore the resulting directory structure to examine how partitioning was applied—each partition key should create a subfolder representing a unique value.
4. Update your data loading, filtering, and aggregation functions (using either **Arrow** or **DuckDB**) to operate on the partitioned **Parquet** file. Then compare the performance with the non-partitioned version.

```

import pyarrow.parquet as pq
dataset = ds.dataset(
 "data/RPindividus.parquet", format="parquet"
).to_table()

pq.write_to_dataset(
 dataset,
 root_path="data/RPindividus_partitionne",
 partition_cols=["REGION", "DEPT"]
)

```



```

import pyarrow.dataset as ds
import pyarrow.compute as pc
import duckdb

@measure_performance
def summarize_filter_parquet_partitioned_arrow(*args, **kwargs):

 dataset = ds.dataset("data/RPindividus_partitionne/", partitioning="hive")
 table = dataset.to_table(filter=pc.field("DEPT").isin(['18', '28', '36']))

 grouped_table = (
 table
 .group_by(["AGED", "DEPT"])
 .aggregate(["IPONDI", "sum"])
 .rename_columns(["AGED", "DEPT", "n_indiv"])
 .to_pandas()
)

 return (
 grouped_table
)

@measure_performance
def summarize_filter_parquet_complete_arrow(*args, **kwargs):

 dataset = ds.dataset("data/RPindividus.parquet")
 table = dataset.to_table(filter=pc.field("DEPT").isin(['18', '28', '36']))

 grouped_table = (
 table
 .group_by(["AGED", "DEPT"])
 .aggregate(["IPONDI", "sum"])
 .rename_columns(["AGED", "DEPT", "n_indiv"])
 .to_pandas()
)

 return (
 grouped_table
)

@measure_performance
def summarize_filter_parquet_complete_duckdb(*args, **kwargs):
 con = duckdb.connect(":memory:")
 query = """
 FROM read_parquet('data/RPindividus.parquet')
 SELECT AGED, DEPT, SUM(IPONDI) AS n_indiv
 WHERE DEPT IN ('11','31','34')
 GROUP BY AGED, DEPT
 """

 return (con.sql(query).to_df())

@measure_performance
def summarize_filter_parquet_partitioned_duckdb(*args, **kwargs):
 con = duckdb.connect(":memory:")

```

```
query = """
FROM read_parquet('data/RPindividus_partitionne/**/*.parquet', hive_partitioning = True)
 SELECT AGED, DEPT, SUM(IPONDI) AS n_indiv
 WHERE DEPT IN ('11','31','34')
 GROUP BY AGED, DEPT
"""

return (con.sql(query).to_df())

summarize_filter_parquet_complete_arrow()
summarize_filter_parquet_partitioned_arrow()
summarize_filter_parquet_complete_duckdb()
summarize_filter_parquet_partitioned_duckdb()
```

? When delivering data in **Parquet** format, how should you choose the partitioning key(s)? What limitations should you keep in mind?

## 3 Data in the Cloud

Cloud storage, in the context of data science, follows the same principle as services like **Dropbox** or **Google Drive**: users can access remote files as if they were stored locally on their machines<sup>4</sup>. In other words, for a **Python** user, **working with cloud-stored files can feel exactly the same** as working with local files.

However, unlike a local path such as **My Documents/mysuperfile**, the files are not physically stored on the user's computer. They are hosted on a remote server, and every operation—reading or writing—relies on a network connection.

### 3.1 Why Not Use **Dropbox** or **Google Drive**?

Despite their similarities, services like **Dropbox** or **Google Drive** are not intended for large-scale data storage and processing. For data-intensive use cases, it is strongly recommended to rely on dedicated storage technologies (see [the production deployment course](#)).

All major cloud providers—AWS, Google Cloud Platform (GCP), Microsoft Azure—rely on a common principle: **object storage**, often implemented through **S3**-compatible systems.

This is why leading cloud platforms offer specialized storage services, typically based on object storage, with **S3** (Simple Storage Service) being the most widely used and recognized standard.

### 3.2 The **S3** System

**S3** (Simple Storage Service), developed by Amazon, has become the de facto standard for cloud-based storage. It is:

- **Reliable**: Data is replicated across multiple servers or zones;
- **Secure**: It supports encryption and fine-grained access control;

---

<sup>4</sup>This functionality is often enabled through virtual file systems or filesystem wrappers compatible with libraries such as **pandas**, **pyarrow**, **duckdb**, and others.

- **Scalable:** It is designed to handle massive volumes of data without performance degradation.

### 3.2.a The Concept of a Bucket

The fundamental unit in **S3** is the **bucket**—a storage container (either private or public) that can contain a virtual file system of folders and files.

To access a file stored in a bucket:

- The user must be **authorized**, typically via credentials or secure access tokens;
- Once authenticated, the user can **read, write, or modify** the contents of the bucket, similar to interacting with a remote file system.

#### **i** Note

The following examples are fully reproducible for SSP Cloud platform users.

They can also be used by AWS users—simply update the **endpoint** URL as shown below.

## 3.3 How to do it with Python?

### 3.3.a Key libraries

Interaction between a remote file system and a local Python session is made possible through APIs. The two main libraries for this purpose are:

- **boto3**, a library developed and maintained by Amazon;
- **s3fs**, a library that enables interaction with stored files as if they were part of a traditional local filesystem.

The **pyarrow** and **duckdb** libraries we previously introduced also support working with cloud-stored data as if it were located on the local machine. This functionality is extremely convenient and helps ensure reliable reading and writing of files in cloud-based environments.

#### **i** Note

On SSP Cloud, access tokens for S3 storage are automatically injected into services when they are launched. These tokens remain valid for 7 days.

If the service icon changes from green to red, it indicates that the tokens have expired — you should save your code and data, then restart the session by launching a new service.

## 3.4 Practical Case: Storing Your Project’s Data on SSP Cloud

A key criterion for evaluating **Python** projects is **reproducibility**—the ability to obtain the same results using the same input data and code. Whenever possible, your final submission should begin with the raw data used as input for your project. If the source files are publicly available via a URL, the ideal approach is to import them directly at the start of your project (see the [Pandas lab](#) for an example of such an import using **Pandas**).

In practice, this is not always feasible. Your data may not be publicly accessible, or it might come in complex formats that require preprocessing before it can be used in standard workflows. In other cases, your dataset may be the result of an automated retrieval process—such as through an **API** or **web scraping**—which can be time-consuming to reproduce. Moreover, since websites frequently change over time, it is often better to “freeze” the data once collected. Similarly, even if you are

not storing raw data, you may want to preserve trained models, as the training process can also be resource-intensive.

In all of these situations, you need a reliable way to store data or models. **Your `Git` repository is not the appropriate place to store large files.** A well-structured `Python` project follows a modular design: it separates code (versioned with `Git`), configuration elements (such as API tokens, which should never be hardcoded), and data storage. This conceptual separation leads to cleaner, more maintainable projects.

While `Git` is designed for source code management, file storage requires dedicated solutions. Many such tools exist. On SSP Cloud, the recommended option is `MinIO`, an open-source implementation of the `S3` storage protocol introduced earlier. This brief tutorial will walk you through a standard workflow for using `MinIO` in your project.

#### ⚠ Warning

No matter which storage solution you use for your data or models, **\*\*you must include the code that generates these objects in your project repository**

### 3.4.a Sharing files on SSP Cloud

As mentioned earlier, files are stored on `S3` in a bucket. On SSP Cloud, a bucket is automatically created when your account is set up, and it has the same name as your SSP Cloud username. The `My Files` interface allows you to access it visually, upload files, download them, and more.

In this tutorial, however, we'll access it programmatically using `Python` code. The `s3fs` package lets you interact with your bucket just like a regular file system. For example, you can list the files available in your *bucket* using the following command:

```
import s3fs

fs = s3fs.S3FileSystem(client_kwargs={"endpoint_url": "https://minio.lab.sspcloud.fr"})

MY_BUCKET = "mon_nom_utilisateur_sspcloud"
fs.ls(MY_BUCKET)
```

If you've never added any files to `MinIO`, your *bucket* will be empty, so the command above should return an empty list. Let's add a first folder to see the difference.

By default, a bucket is personal, meaning the data inside can only be read or modified by you. For your project, you'll want to share these files with your group members to collaborate efficiently. But that's not all! Your instructors will also need to access these files to reproduce your analyses.

There are several ways to make files more or less public on `MinIO`. The simplest—and the one we recommend—is to create a folder named `diffusion` at the root of your bucket. On SSP Cloud, all files inside a `diffusion` folder are **readable** by all authenticated users. Use the `My Files` interface to create a `diffusion` folder at the root of your *bucket*. If everything worked correctly, the `Python` command above should now display the path `your_sspcloud_username/diffusion`.

#### i Cloud storage encourages collaborative work!

Instead of each project member working with their own files on their personal machine—which would require frequent syncing and makes reproducibility more error-prone—files are placed in a central repository that every group member can access.

To do this, the group should agree to use one member's bucket, and ensure that other members can access the data by placing it in the `diffusion` folder of the selected bucket.

### 3.4.b Retrieving and storing data

Now that we know where to place our data on **MinIO**, let's look at how to do it in practice using **Python**.

#### Case of a DataFrame

Let's revisit an example from the course on **APIs** to simulate a time-consuming data retrieval step.

```
import requests
import pandas as pd

url_api = "https://koumoul.com/data-fair/api/v1/datasets/dpe-france/lines?format=json&q_mode=simple&q_s=code_insee_commune_actualise%3A%2201450%22&size=100&select=%2A&sampling=neighbors"
response_json = requests.get(url_api).json()
df_dpe = pd.json_normalize(response_json["results"])

df_dpe.head(2)
```

class_communication_energie	trfct_mobilier_dpe_type_batelle	insee_construction	propret	latitude	surface_batiment_sqm	numero_dpe	j	trfct_type_batiment_description	geo_adresse	geo_ville	class_estimation_geo	nom_methode_dpe	trfct_dapartement_cade	communication_energie	date_etablissement_dpe	longitude	score	id	version_methode_dpe
0	E	Vente	1	45.9274843.233095	45.927488	106.87	1301Y20000015	2	Maison Individuelle	Rue du Château 01000 Villars-Loyse-Molien	0.50	0	Méthode Facture	01	2013-04-15	5.200195	None	10145014507014502204Q404	0401
1	G	Vente	1000	45.9313763.230461	45.931376	70.78	1301Y20000018	9	Maison Individuelle	552 Rue Royale 01000 Villars-Loyse-Molien	0.54	15	Méthode ICL	01	2013-04-22	5.200461	None	10145014507014502204Q404	0401

This request returns a **Pandas DataFrame**, and the first two rows are printed above. In our example, the process is deliberately simple, but in practice, you might have many steps of querying and preparing the data before obtaining a usable DataFrame for the rest of the project. This process might be time-consuming, so we'll store these “intermediate” data on **MinIO** to avoid rerunning all the code that generated them each time.

We can use **Pandas** export functions, which allow saving data in various formats. Since we're working in the cloud, one additional step is needed: we open a connection to **MinIO**, then export our DataFrame.

```
MY_BUCKET = "mon_nom_utilisateur_sspcloud"
FILE_PATH_OUT_S3 = f"{MY_BUCKET}/diffusion/df_dpe.csv"

with fs.open(FILE_PATH_OUT_S3, 'w') as file_out:
 df_dpe.to_csv(file_out)
```

You can verify that your file has been successfully uploaded either via the **My Files** interface, or directly in **Python** by checking the contents of the **diffusion** folder in your **bucket**:

```
fs.ls(f"{MY_BUCKET}/diffusion")
```

We could just as easily export our dataset in **Parquet** format to reduce storage space and improve read performance. Note: since **Parquet** is a compressed format, you must specify that you're writing a binary file — the file opening mode passed to **fs.open** should be changed from **w** (write) to **wb** (write binary).

```
FILE_PATH_OUT_S3 = f"{MY_BUCKET}/diffusion/df_dpe.parquet"

with fs.open(FILE_PATH_OUT_S3, 'wb') as file_out:
 df_dpe.to_parquet(file_out)
```

#### File-based use case

In the previous section, we dealt with the “simple” case of a DataFrame, which allowed us to use **Pandas** built-in export functions. Now, let's imagine we have multiple input files, each potentially in a different format. A typical example of such files are **ShapeFiles**, which are geographic data files and are composed of several related files (see **GeoPandas chapter**). Let's start by downloading a **.shp** file to inspect its structure.

Below are the outlines of the department of Réunion produced by the IGN, in the form of a `..7z` archive that we will unzip locally into a folder called `departements_fr`.

```
import io
import os
import requests
import py7zr

Import et décompression
contours_url = "https://data.geopf.fr/telechargement/download/ADMIN-EXPRESS-COG-CARTO/ADMIN-EXPRESS-COG-CARTO_3-2__SHP_RGR92UTM40S_REU_2023-05-03/ADMIN-EXPRESS-COG-CARTO_3-2__SHP_RGR92UTM40S_REU_2023-05-03.7z"
response = requests.get(contours_url, stream=True)

with py7zr.SevenZipFile(io.BytesIO(response.content), mode='r') as archive:
 archive.extractall(path="departements_fr")

Vérification du dossier (local, pas sur S3)
os.listdir("departements_fr")
```

```
['ADMIN-EXPRESS-COG-CARTO_3-2__SHP_RGR92UTM40S_REU_2023-05-03']
```

Since we're now dealing with local files rather than a `Pandas DataFrame`, we need to use the `s3fs` package to transfer files from the local filesystem to the remote filesystem (`MinIO`). Using the `put` command, we can copy an entire folder to `MinIO` in a single step. Be sure to set the `recursive=True` parameter so that both the folder and its contents are copied.

```
fs.put("departements_fr/", f"{MY_BUCKET}/diffusion/departements_fr/", recursive=True)
```

Let's check that the folder was successfully copied:

```
fs.ls(f"{MY_BUCKET}/diffusion/departements_fr")
```

If the folder, as in this case, contains files at multiple levels, you will need to browse the list recursively to access the files. If, for example, you are only interested in municipal divisions, you can use a `glob`:

```
fs.glob(f"{MY_BUCKET}/diffusion/departements_fr/**/COMMUNE.*')
```

If everything worked correctly, the command above should return a list of file paths on `MinIO` for the various components of the Réunion cities `ShapeFile` (`.shp`, `.shx`, `.prj`, etc.).

### 3.4.c Using the data

In the reverse direction, to retrieve files from `MinIO` in a `Python` session, the commands are symmetrical.

#### Case of a dataframe

Make sure to pass the `r` parameter (`read`, for reading) instead of `w` (`write`, for writing) to the `fs.open` function to avoid overwriting the file!

```
MY_BUCKET = "your_sspcloud_username"
FILE_PATH_S3 = f"{MY_BUCKET}/diffusion/df_dpe.csv"

Import
with fs.open(FILE_PATH_S3, 'r') as file_in:
 df_dpe = pd.read_csv(file_in)
```

```
Check
df_dpe.head(2)
```

Similarly, if the file is in **Parquet** format (don't forget to use **rb** instead of **r** due to compression):

```
MY_BUCKET = "your_sspcloud_username"
FILE_PATH_S3 = f"{MY_BUCKET}/diffusion/df_dpe.parquet"

Import
with fs.open(FILE_PATH_S3, 'rb') as file_in:
 df_dpe = pd.read_parquet(file_in)

Check
df_dpe.head(2)
```

### Case of files

For file collections, you'll first need to download the files from **MinIO** to your local machine (i.e., the current SSP Cloud session).

```
Retrieve files from MinIO to the local machine
fs.get(f"{MY_BUCKET}/diffusion/departements_fr/", "departements_fr/", recursive=True)
```

Then, you can import them in the usual way using the appropriate **Python** package. For **ShapeFiles**, where multiple files make up a single dataset, one command is sufficient after retrieval:

```
from pathlib import Path
import geopandas as gpd

localpath = Path("departements_fr")

df_dep = gpd.read_file(
 list(
 localpath.glob("**/COMMUNE.shp")
)[0]
)
df_dep.head(2)
```

#### Line 4

As our file contains many files, we will need to search through it using, once again, a **glob**

#### Line 10

**PathLib** gives us a list, so we restrict ourselves to the first echo and import it with **GeoPandas**

## 3.5 To go further

- [SSPCloud documentation on MinIO](#)

## Informations additionnelles

**i** Python environment

This site was built automatically through a [Github](#) action using the [Quarto](#) reproducible publishing software (version 1.10.18).

The environment used to obtain the results is reproducible via [uv](#). The [pyproject.toml](#) file used to build this environment is available on the [linogaliana/python-datascientist](#) repository

```
[project]
name = "python-datascientist"
version = "0.1.0"
description = "Source code for Lino Galiana's Python for data science course"
readme = "README.md"
requires-python = "≥3.13,<3.14"
dependencies = [
 "altair≥6.0.0",
 "cartiflette",
 "contextily=1.6.2",
 "duckdb≥0.10.1",
 "folium≥0.19.6",
 "gdal=3.11.4",
 "graphviz=0.20.3",
 "great-tables≥0.12.0",
 "gt-extras≥0.0.8",
 "ipykernel≥6.29.5",
 "jupyter≥1.1.1",
 "jupyter-cache≥1.0.0",
 "kaleido≥0.2.1",
 "langchain-community≥0.3.27",
 "loguru=0.7.3",
 "markdown≥3.8",
 "nbcClient≥0.10.0",
 "nbformat≥5.10.4",
 "nltk≥3.9.1",
 "pandas≥3.0",
 "pip≥25.1.1",
 "plotly≥6.1.2",
 "plotnine≥0.15",
 "polars≥1.8.2",
 "pyarrow≥17.0.0",
 "pynsee≥0.1.8",
 "python-dotenv≥1.0.1",
 "python-frontmatter≥1.1.0",
 "pywaffle≥1.1.1",
 "requests≥2.32.3",
 "scikit-image≥0.24.0",
 "scikit-learn≥1.8.0",
 "scipy≥1.13.0",
 "seaborn≥0.13.2",
 "selenium<4.39.0",
 "spacy≥3.8.4",
 "webdriver-manager≥4.0.2",
 "wordcloud=1.9.3",
]

[tool.uv.sources]
cartiflette = { git = "https://github.com/insee-frlab/cartiflette" }
gdal = [
 { index = "gdal-wheels", marker = "sys_platform == 'linux'", },
 { index = "geospatial-wheels", marker = "sys_platform == 'win32'", },
]
```



```
]

[[tool.uv.index]]
name = "geospatial_wheels"
url = "https://nathanjmcougall.github.io/geospatial-wheels-index/"
explicit = true

[[tool.uv.index]]
name = "gdal-wheels"
url = "https://gitlab.com/api/v4/projects/61637378/packages/pypi/simple"
explicit = true

[dependency-groups]
dev = [
 "nb-clean ≥ 4.0.1",
]
```

To use exactly the same environment (version of `Python` and `packages`), please refer to the [documentation for uv](#).

## Bibliography

- [1] A. Dondon and P. Lamarche, “Quels formats pour quelles données?” *Courrier des statistiques*, pp. 86–103, 2023.